




HANDSON 2

Big Data Systems

Soheil Shirvani - 3720505




Contents

Lunching Stado Parallel database	2
Steps for lunching Stado Parallel:	2
Task 1	4
1.a Create a partitioned table partab	4
1.b Insert the tuples (with given values) into partab	5
2.a Create a replicated table reptab.	6
2.b Insert the tuples (with given values) into reptab	7
Results Only:	8
Schema of your table partab (output of: show table partab)	8
Output of the query: SELECT * FROM partab	8
Schema of your table reptab (output of: show table reptab)	8
Output of the query: SELECT * FROM reptab	8
Task 2	9
Run 3 query on Stado as parallel database	9
Query 1:	9
Query 3:	10
Query 6:	11
Run 3 query on PostgreSQL as Single database.	12
Query 1:	12
Query 3:	13
Query 6:	14
Results Only	15
Average execution time of the warm runs of Q1, Q3 and Q6 with Stado	15
Average execution time of the warm runs of Q1, Q3 and Q6 with Single DB	15
Speedup of Q1, Q3 and Q6 achieved with Stado against single instance PostgreSQL	15


Lunching Stado Parallel database

Steps for lunching Stado Parallel:


1. Here first we login into 4 different VMs. 1 is master VM and 3 Slave VMs. To do so we open all the VMs and login.



2. Start single PostgreSQL database on all the VMs. **\$ pg_ctl start**



3. On master VM we will start parallel Stado database. **\$ gs-server.sh -d testdb**



4. We will start Stado SQL client to connect to the parallel database. **\$ gs-cmdline.sh -d testdb -u admin -p admin -z**

```

fcs-vm-bigdata-master [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places
bigdata@bigdatamaster: ~/stado/stado/bin
File Edit View Search Terminal Help
bigdata@bigdatamaster:~$ pg_ctl start
pg_ctl: another server might be running; trying to start server anyway
server starting
bigdata@bigdatamaster:~$ LOG: database system was interrupted; last known up at
2017-02-03 12:17:01 AST
LOG: database system was not properly shut down; automatic recovery in progress
LOG: record with zero length at 2/81847D18
LOG: redo is not required
LOG: database system is ready to accept connections
LOG: autovacuum launcher started

bigdata@bigdatamaster:~$ cd /home/bigdata/stado/stado/bin/
bigdata@bigdatamaster:~/stado/stado/bin$ gs-server.sh -d testdb
Starting....
bigdata@bigdatamaster:~/stado/stado/bin$ gs-cmdline.sh -d testdb -u admin -p adm
in -z

Stado ->

```

5. Then we are able to see the tables in the database or create and insert in it.

```

fcs-vm-bigdata-master [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places
bigdata@bigdatamaster: ~/stado/stado/bin
File Edit View Search Terminal Help
bigdata@bigdatamaster:~$ cd /home/bigdata/stado/stado/bin/
bigdata@bigdatamaster:~/stado/stado/bin$ gs-server.sh -d testdb
Starting....
bigdata@bigdatamaster:~/stado/stado/bin$ gs-cmdline.sh -d testdb -u admin -p adm
in -z

Stado -> show tables;
+-----+
| TABLE | TABLE_PARTITIONING_COLUMN | TABLE_NODES |
+-----+
| customer | c_custkey | 1,2,3,4 |
| lineitem | l_orderkey | 1,2,3,4 |
| nation | n_nationkey | 1,2,3,4 |
| orders | o_orderkey | 1,2,3,4 |
| part | p_partkey | 1,2,3,4 |
| partsupp | ps_partkey | 1,2,3,4 |
| region | r_regionkey | 1,2,3,4 |
| supplier | s_suppkey | 1,2,3,4 |
+-----+
8 row(s).
Response time: 0m 0s 159ms Total time: 0m 0s 164ms

Stado ->

```

Task 1

1.a Create a partitioned table partab

To Create a partitioned table named partab with 2 column including (col1, Integer) and (col2, Char) and partition key of (col1) we use the query below:

```
CREATE TABLE partab (col1 INTEGER, col2 CHAR(15)) PARTITIONING KEY col1 ON ALL;
```

And response is:

```
Stado -> CREATE TABLE partab (col1 INTEGER, col CHAR(15)) PARTITIONING KEY col1
ON ALL;
OK
Response time: 0m 1s 275ms Total time: 0m 1s 276ms
```

So the table is created. And Show tables here include the new table:

```
Stado -> show tables;
+-----+-----+-----+
| TABLE | TABLE_PARTITIONING_COLUMN | TABLE_NODES |
+-----+-----+-----+
| customer | c_custkey | 1,2,3,4 |
| lineitem | l_orderkey | 1,2,3,4 |
| nation | | 1,2,3,4 |
| orders | o_orderkey | 1,2,3,4 |
| part | p_partkey | 1,2,3,4 |
| partab | col1 | 1,2,3,4 |
| partsupp | ps_partkey | 1,2,3,4 |
| region | | 1,2,3,4 |
| supplier | | 1,2,3,4 |
+-----+-----+-----+
9 row(s).
Response time: 0m 0s 5ms Total time: 0m 0s 7ms
```

1.b Insert the tuples (with given values) into partab

We will Insert 6 predefined tuples into our partab table that we created before. We use the command Below:

INSERT INTO partab VALUES ('2','I'), ...

```

bigdata@bigdatamaster: ~/stado/stado/bin
File Edit View Search Terminal Help
1 row(s) affected
Response time: 0m 0s 7ms Total time: 0m 0s 7ms

Stado -> INSERT INTO partab VALUES (3,'B');
1 row(s) affected
Response time: 0m 0s 5ms Total time: 0m 0s 5ms

Stado -> INSERT INTO partab VALUES (4,'G');
1 row(s) affected
Response time: 0m 0s 7ms Total time: 0m 0s 7ms

Stado -> INSERT INTO partab VALUES (5,'Y');
1 row(s) affected
Response time: 0m 0s 8ms Total time: 0m 0s 8ms

Stado -> INSERT INTO partab VALUES (6,'O');
1 row(s) affected
Response time: 0m 0s 4ms Total time: 0m 0s 4ms

Stado -> INSERT INTO partab VALUES (7,'R');
1 row(s) affected
Response time: 0m 0s 5ms Total time: 0m 0s 5ms

Stado ->

```

After Inserting all the tuples we can query our table and check if the tuples are in the table using the command below and see the result:

*SELECT * FROM partab*

```

Stado -> SELECT * FROM partab;
+-----+
| col1 | col |
+-----+
| 1 | V |
| 2 | I |
| 7 | R |
| 3 | B |
| 4 | G |
| 7 | R |
| 5 | Y |
| 6 | O |
+-----+
8 row(s).
Response time: 0m 0s 49ms Total time: 0m 0s 50ms

```

Table Actually have all the rows (7,'R') is not displayed and it was inserted twice. So all the insertion worked perfectly fine.

2.a Create a replicated table reptab.

To Create a replicated table named reptab with 2 column including (fld1, Integer) and (fld2, Char(2)) we use the query below:

```
CREATE TABLE reptab (fld1 INTEGER, fld2 CHAR(2)) REPLICATED;
```

And response is:

```
Stado -> CREATE TABLE reptab (fld1 INTEGER, fld2 CHAR(2)) REPLICATED;
OK
Response time: 0m 0s 106ms Total time: 0m 0s 106ms
```

So the table is created. And Show tables here include the new table:

```
Stado -> show tables;
```

TABLE	TABLE_PARTITIONING_COLUMN	TABLE_NODES
customer	c_custkey	1,2,3,4
lineitem	l_orderkey	1,2,3,4
nation		1,2,3,4
orders	o_orderkey	1,2,3,4
part	p_partkey	1,2,3,4
partab	col1	1,2,3,4
partsupp	ps_partkey	1,2,3,4
region		1,2,3,4
reptab		1,2,3,4
supplier		1,2,3,4

```
10 row(s).
Response time: 0m 0s 16ms Total time: 0m 0s 17ms
```


2.b Insert the tuples (with given values) into reptab

We will Insert 6 predefined tuples into our reptab table that we created before. We use the command Below:

INSERT INTO reptab VALUES ('2','I'), ...

```
Stado -> INSERT INTO reptab VALUES (7,'NB');
1 row(s) affected
Response time: 0m 0s 18ms Total time: 0m 0s 19ms

Stado -> INSERT INTO reptab VALUES (8,'NL');
1 row(s) affected
Response time: 0m 0s 12ms Total time: 0m 0s 12ms

Stado -> INSERT INTO reptab VALUES (9,'NS');
1 row(s) affected
Response time: 0m 0s 13ms Total time: 0m 0s 15ms

Stado -> INSERT INTO reptab VALUES (10, 'PE');
1 row(s) affected
Response time: 0m 0s 11ms Total time: 0m 0s 11ms

Stado -> 
```

After Inserting all the tuples we can query our table and check if the tuples are in the table using the command below and see the result:

*SELECT * FROM reptab*

```
Stado -> SELECT * FROM reptab;
+-----+
| fld1 | fld2 |
+-----+
|    7 | NB   |
|    8 | NL   |
|    9 | NS   |
|   10 | PE   |
+-----+
4 row(s).
Response time: 0m 0s 11ms Total time: 0m 0s 12ms

Stado ->
```

As we can see all the tuples were inserted into the reptab successfully.

Results Only:

1. Partitioned

CREATE TABLE partab (col1 INTEGER, col2 CHAR) PARTITIONING KEY col1 ON ALL;

Schema of your table partab (output of: show table partab)

```

+-----+
| COLUMN_NAME | SQL_DATA_TYPE | TYPE_NAME | IS_NULLABLE | KEY | DEFAULT |
+-----+
| col1        | 4             | INTEGER   | YES         | NO  |         |
| col         | 1             | CHAR(15)  | YES         | NO  |         |
+-----+
2 row(s).
Response time: 0m 0s 5ms  Total time: 0m 0s 5ms

```

Output of the query: SELECT * FROM partab

```

Stado -> SELECT * FROM partab;
+-----+
| col1 | col |
+-----+
| 1    | V   |
| 2    | I   |
| 7    | R   |
| 3    | B   |
| 4    | G   |
| 7    | R   |
| 5    | Y   |
| 6    | O   |
+-----+
8 row(s).
Response time: 0m 0s 49ms  Total time: 0m 0s 50ms

```

2. Replicated

Schema of your table reptab (output of: show table reptab)

```

Stado -> show table reptab
Stado -> ;
+-----+
| COLUMN_NAME | SQL_DATA_TYPE | TYPE_NAME | IS_NULLABLE | KEY | DEFAULT |
+-----+
| fld1        | 4             | INTEGER   | YES         | NO  |         |
| fld2        | 1             | CHAR(2)   | YES         | NO  |         |
+-----+
2 row(s).
Response time: 0m 0s 3ms  Total time: 0m 0s 4ms

```

Output of the query: SELECT * FROM reptab

```

Stado -> SELECT * FROM reptab;
+-----+
| fld1 | fld2 |
+-----+
| 7    | NB   |
| 8    | NL   |
| 9    | NS   |
| 10   | PE   |
+-----+
4 row(s).
Response time: 0m 0s 11ms  Total time: 0m 0s 12ms
Stado ->

```

Task 2

Here we are going to run 3 queries from TPC-H benchmark and find the response time of database in 2 modes Parallel and Single. I ignored the first run of each query (cold) and made a note of the response time of 2 next query execution for both modes.

Run 3 query on Stado as parallel database.

Query 1:

After running 1 time (cold run) the Second result is: **Time: 13s 689ms**

```
bigdata@bigdatamaster: ~/stado/stado/bin x bigdata@bigdatamaster: ~/stado/stado/bin x
Stado -> select l_returnflag, l_linestatus, sum(l_quantity) as sum_qty, sum(l_extendedprice) as sum_base_price, sum(l_extendedprice * (1-l_discount)) as sum_disc_price,
sum(l_extendedprice * (1-l_discount) * (1+l_tax)) as sum_charge, avg(l_quantity) as avg_qty, avg(l_extendedprice) as avg_price, avg(l_discount) as avg_disc, count(*) as
count_order from linetitem where l_shipdate <= (date '1998-12-01' - interval '110 days') group by l_returnflag, l_linestatus order by l_returnflag, l_linestatus;
WARNING: there is already a transaction in progress
WARNING: there is already a transaction in progress
WARNING: there is already a transaction in progress
WARNING: there is already a transaction in progress
+-----+
| l_returnflag | l_linestatus | sum_qty | sum_base_price | sum_disc_price | sum_charge | avg_qty | avg_price | avg_disc |
+-----+
| A | F | 75478173.00 | 113197331346.02 | 107536408207.3092 | 111838898769.61761 | 25.5056988774585058 | 38251.814164122399 | 0.05000395030255803 |
211 | 2959267 |
| N | F | 1966480.00 | 2946114826.74 | 2798796636.1563997 | 2911030163.0685883 | 25.5300807519538857 | 38248.316500142809 | 0.04999558590605769 |
480 | 77026 |
| N | O | 146534288.00 | 219743580642.00 | 208760665820.8363 | 21711398915.43436 | 25.4959210851865604 | 38233.815904061724 | 0.04997896426221282 |
042 | 5747362 |
| R | F | 75577628.00 | 113351914218.17 | 107688081811.48871 | 111994307866.22044 | 25.5121503826098757 | 38263.321543918361 | 0.04997978002421671 |
223 | 2962417 |
+-----+
4 row(s).
Response time: 0m 13s 687ms Total time: 0m 13s 689ms
Stado -> █
```

Third run is: **Time: 13s 961ms**

```
bigdata@bigdatamaster: ~/stado/stado/bin x bigdata@bigdatamaster: ~/stado/stado/bin x
Stado -> select l_returnflag, l_linestatus, sum(l_quantity) as sum_qty, sum(l_extendedprice) as sum_base_price, sum(l_extendedprice * (1-l_discount)) as sum_disc_price,
sum(l_extendedprice * (1-l_discount) * (1+l_tax)) as sum_charge, avg(l_quantity) as avg_qty, avg(l_extendedprice) as avg_price, avg(l_discount) as avg_disc, count(*) as
count_order from linetitem where l_shipdate <= (date '1998-12-01' - interval '110 days') group by l_returnflag, l_linestatus order by l_returnflag, l_linestatus;
WARNING: there is already a transaction in progress
WARNING: there is already a transaction in progress
WARNING: there is already a transaction in progress
WARNING: there is already a transaction in progress
+-----+
| l_returnflag | l_linestatus | sum_qty | sum_base_price | sum_disc_price | sum_charge | avg_qty | avg_price | avg_disc |
+-----+
| A | F | 75478173.00 | 113197331346.02 | 107536408207.3092 | 111838898769.61761 | 25.5056988774585058 | 38251.814164122399 | 0.05000395030255803 |
211 | 2959267 |
| N | F | 1966480.00 | 2946114826.74 | 2798796636.1563997 | 2911030163.0685883 | 25.5300807519538857 | 38248.316500142809 | 0.04999558590605769 |
480 | 77026 |
| N | O | 146534288.00 | 219743580642.00 | 208760665820.8363 | 21711398915.43436 | 25.4959210851865604 | 38233.815904061724 | 0.04997896426221282 |
042 | 5747362 |
| R | F | 75577628.00 | 113351914218.17 | 107688081811.48871 | 111994307866.22044 | 25.5121503826098757 | 38263.321543918361 | 0.04997978002421671 |
223 | 2962417 |
+-----+
4 row(s).
Response time: 0m 13s 945ms Total time: 0m 13s 961ms
Stado -> █
```

Query 3:

After running 1 time (cold run) the Second result is: **Time: 10s 451ms**

```
Stado -> SELECT L_ORDERKEY, SUM(L_EXTENDEDPRICE*(1-L_DISCOUNT)) AS REVENUE, O_ORDERDATE, O_SHIPPRIORITY FROM CUSTOMER, ORDERS, LINEITEM WHERE C_MKTSEGMENT IN ('AUTOMOBILE') AND C_CUSTKEY = O_CUSTKEY AND L_ORDERKEY = O_ORDERKEY AND O_ORDERDATE < date '1995-03-19' AND L_SHIPDATE > date '1995-03-19' GROUP BY L_ORDERKEY, O_ORDERDATE, O_SHIPPRIORITY ORDER BY REVENUE DESC, O_ORDERDATE;
```

```

3389800 | 1095.4752 | 1994-12-11 | 0 |
4302085 | 1091.9898 | 1995-02-22 | 0 |
3452033 | 1076.0937 | 1994-12-01 | 0 |
7557798 | 1066.9282 | 1995-01-28 | 0 |
9988964 | 1066.7136 | 1995-01-01 | 0 |
442437 | 1059.9072 | 1994-12-01 | 0 |
6456420 | 1059.9072 | 1995-02-19 | 0 |
8080036 | 1050.6652 | 1995-03-12 | 0 |
3445601 | 1044.6432 | 1995-01-06 | 0 |
6834311 | 1037.9440 | 1995-03-06 | 0 |
10804000 | 1029.6891 | 1995-01-20 | 0 |
5012578 | 989.7888 | 1994-12-18 | 0 |
7674112 | 977.7034 | 1995-02-02 | 0 |
5613284 | 963.3096 | 1994-12-30 | 0 |
4182179 | 924.4992 | 1994-11-25 | 0 |
7670117 | 915.8496 | 1995-02-27 | 0 |
2204743 | 906.7872 | 1995-03-15 | 0 |
7231653 | 900.9637 | 1994-12-21 | 0 |
9097634 | 883.9502 | 1995-01-14 | 0 |
9418727 | 873.9905 | 1994-12-01 | 0 |
+-----+
22451 row(s).
Response time: 0m 7s 734ms Total time: 0m 10s 451ms

```

Third run is: **Time: 10s 459ms**

```

9988964 | 1066.7136 | 1995-01-01 | 0 |
442437 | 1059.9072 | 1994-12-01 | 0 |
6456420 | 1059.9072 | 1995-02-19 | 0 |
8080036 | 1050.6652 | 1995-03-12 | 0 |
3445601 | 1044.6432 | 1995-01-06 | 0 |
6834311 | 1037.9440 | 1995-03-06 | 0 |
10804000 | 1029.6891 | 1995-01-20 | 0 |
5012578 | 989.7888 | 1994-12-18 | 0 |
7674112 | 977.7034 | 1995-02-02 | 0 |
5613284 | 963.3096 | 1994-12-30 | 0 |
4182179 | 924.4992 | 1994-11-25 | 0 |
7670117 | 915.8496 | 1995-02-27 | 0 |
2204743 | 906.7872 | 1995-03-15 | 0 |
7231653 | 900.9637 | 1994-12-21 | 0 |
9097634 | 883.9502 | 1995-01-14 | 0 |
9418727 | 873.9905 | 1994-12-01 | 0 |
+-----+
22451 row(s).
Response time: 0m 7s 815ms Total time: 0m 10s 459ms

```

Query 6:

After running 1 time (cold run) the Second result is: **Time: 2s 219ms**

```
Stado -> SELECT SUM(L_EXTENDEDPRICE*L_DISCOUNT) AS REVENUE FROM LINEITEM WHERE L_SHIPDATE >= date '1995-01-01' AND L_SHIPDATE < (date '1995-01-01' + interval '1 year') AND L_DISCOUNT BETWEEN .04 - 0.01 AND .04 + 0.01 AND L_QUANTITY < 24;
```

```
+-----+
|  revenue  |
+-----+
| 164347732.3953 |
+-----+
1 row(s).
Response time: 0m 2s 219ms  Total time: 0m 2s 219ms
```

Third run is: **Time: 1s 909ms**

```
+-----+
|  revenue  |
+-----+
| 164347732.3953 |
+-----+
1 row(s).
Response time: 0m 1s 909ms  Total time: 0m 1s 909ms

Stado -> █
```


Run 3 query on PostgreSQL as Single database.

First we stop the Stado and go to PostgreSQL command line to do the queries.

```
Stado -> ^C
bigdata@bigdatamaster:~/stado/stado/bin$ cd
bigdata@bigdatamaster:~$ psql testdb
psql (9.4.0)
Type "help" for help.

testdb=#
```

Query 1:

After running 1 time (cold run) the Second result is: **Time: 40s 248ms**

```
testdb=# select l_returnflag, l_linestatus, sum(l_quantity) as sum_qty, sum(l_extendedprice) as sum_base_price, sum(l_extendedprice * (1-l_discount)) as sum_disc_price,
sum(l_extendedprice * (1-l_discount) * (1+l_tax)) as sum_charge, avg(l_quantity) as avg_qty, avg(l_extendedprice) as avg_price, avg(l_discount) as avg_disc, count(*) as
count_order from lineitem where l_shipdate <= (date '1998-12-01' - interval '110 days') group by l_returnflag, l_linestatus order by l_returnflag, l_linestatus;
```

l_returnflag	l_linestatus	sum_qty	sum_base_price	sum_disc_price	sum_charge	avg_qty	avg_price	avg_disc
A	F	75478173.00	113197331346.02	107536408207.3092	111838898769.617614	25.5056988774585058	38251.814164122399	0.050003950302558032
N	F	1966480.00	2946114826.74	2798796636.1564	2911030163.068588	25.5300807519538857	38248.316500142809	0.049995585906057694
N	O	146534288.00	219743580642.00	208760665820.8363	217113989915.434356	25.4959210851865604	38233.815904061724	0.049978964262212820
R	F	75577628.00	113351914218.17	107688081811.4887	111994307866.220439	25.5121503826098757	38263.321543918361	0.049979780024216712

(4 rows)

Time: 40248.856 ms

Third run is: **Time: 40s 788ms**

l_returnflag	l_linestatus	sum_qty	sum_base_price	sum_disc_price	sum_charge	avg_qty	avg_price	avg_disc
A	F	75478173.00	113197331346.02	107536408207.3092	111838898769.617614	25.5056988774585058	38251.814164122399	0.050003950302558032
N	F	1966480.00	2946114826.74	2798796636.1564	2911030163.068588	25.5300807519538857	38248.316500142809	0.049995585906057694
N	O	146534288.00	219743580642.00	208760665820.8363	217113989915.434356	25.4959210851865604	38233.815904061724	0.049978964262212820
R	F	75577628.00	113351914218.17	107688081811.4887	111994307866.220439	25.5121503826098757	38263.321543918361	0.049979780024216712

(4 rows)

Time: 40788.093 ms

Query 3:

After running 1 time (cold run) the Second result is: **Time: 16s 948ms**

```
testdb=# SELECT L_ORDERKEY, SUM(L_EXTENDEDPRICE*(1-L_DISCOUNT)) AS REVENUE, O_ORDERDATE, O_SHIPPRIORITY FROM CUSTOMER, ORDERS, LINEITEM WHERE C_MKTSEGMENT IN ('AUTOMOBILE') AND C_CUSTKEY = O_CUSTKEY AND L_ORDERKEY = O_ORDERKEY AND O_ORDERDATE < date '1995-03-19' AND L_SHIPDATE > date '1995-03-19' GROUP BY L_ORDERKEY, O_ORDERDATE, O_SHIPPRIORITY ORDER BY REVENUE DESC, O_ORDERDATE;
```

1824643	380013.3687	1995-03-06	0
1088995	375667.8475	1995-03-11	0
6669831	375492.0278	1995-02-07	0
4710272	375188.6517	1995-03-17	0
3778628	373654.2087	1995-02-25	0
4736678	372807.1864	1995-03-03	0
1062786	372672.5509	1995-02-10	0
6680197	369266.7219	1995-02-13	0
10350439	369187.7220	1995-02-16	0
8257025	368258.2766	1995-03-16	0
4496835	367877.7452	1995-03-18	0
7205063	366929.7907	1995-02-26	0
1102534	366089.2778	1995-03-15	0
7166048	365379.9582	1995-03-02	0

Time: 16948.201 ms

Third run is: **Time: 16s 187ms**

4710272	375188.6517	1995-03-17	0
3778628	373654.2087	1995-02-25	0
4736678	372807.1864	1995-03-03	0
1062786	372672.5509	1995-02-10	0
6680197	369266.7219	1995-02-13	0
10350439	369187.7220	1995-02-16	0
8257025	368258.2766	1995-03-16	0
4496835	367877.7452	1995-03-18	0
7205063	366929.7907	1995-02-26	0
1102534	366089.2778	1995-03-15	0
7166048	365379.9582	1995-03-02	0

Time: 16187.535 ms

Query 6:

After running 1 time (cold run) the Second result is: **Time: 32s 14ms**

```
testdb=# SELECT SUM(L_EXTENDEDPRICE*L_DISCOUNT) AS REVENUE FROM LINEITEM WHERE L_SHIPDATE >= date '1995-01-01' AND L_SHIPDATE < (date '1995-01-01' + interval '1 year') A  
ND L_DISCOUNT BETWEEN .04 - 0.01 AND .04 + 0.01 AND L_QUANTITY < 24;
```

```
      revenue  
-----  
 164347732.3953  
(1 row)  
  
Time: 32014.984 ms  
testdb=#
```

Third run is: **Time: 31s 888ms**

```
      revenue  
-----  
 164347732.3953  
(1 row)  
  
Time: 31888.667 ms  
testdb=#
```

Results Only

Average execution time of the warm runs of Q1, Q3 and Q6 with Stado

$$Q1 \begin{cases} \text{warm run 1 : 13s 689ms : 13689ms} \\ \text{warm run 2 : 13s 961ms : 13961ms} \end{cases} \rightarrow Avg = 13825ms$$

$$Q3 \begin{cases} \text{warm run 1 : 10s 451ms : 10451ms} \\ \text{warm run 2 : 10s 459ms : 10459ms} \end{cases} \rightarrow Avg = 10455ms$$

$$Q6 \begin{cases} \text{warm run 1 : 2s 219ms : 2219ms} \\ \text{warm run 2 : 1s 909ms : 1909ms} \end{cases} \rightarrow Avg = 2064ms$$

Average execution time of the warm runs of Q1, Q3 and Q6 with Single DB

$$Q1 \begin{cases} \text{warm run 1 : 40s 248ms : 40248ms} \\ \text{warm run 2 : 40s 788ms : 40788ms} \end{cases} \rightarrow Avg = 40518ms$$

$$Q3 \begin{cases} \text{warm run 1 : 16s 948ms : 16948ms} \\ \text{warm run 2 : 16s 187ms : 16187ms} \end{cases} \rightarrow Avg = 16568ms$$

$$Q6 \begin{cases} \text{warm run 1 : 32s 14ms : 32014ms} \\ \text{warm run 2 : 31s 888ms : 31888ms} \end{cases} \rightarrow Avg = 31951ms$$

Speedup of Q1, Q3 and Q6 achieved with Stado against single instance
PostgreSQL

$$Q1 \text{ Speedup} = \frac{40518}{13825} = 2.93$$

$$Q3 \text{ Speedup} = \frac{16568}{10455} = 1.584$$

$$Q6 \text{ Speedup} = \frac{31951}{2064} = 15.48$$